

Data Dictionary

Saeid Abolfazli

May 2, 2016

The file “data/climate_change.csv” contains climate data from **May 1983** to **December 2008**. The data and below description are provided by team of [MITx: 15.071x The Analytics Edge](#) @ EDX

Features:

The available variables include:

- **Year:** the observation year.
- **Month:** the observation month.
- **Temp:** the difference in degrees Celsius between the average global temperature in that period and a reference value. This data comes from the Climatic Research Unit at the University of East Anglia.
- **CO2, N2O, CH4, CFC.11, CFC.12:** atmospheric concentrations of carbon dioxide (CO2), nitrous oxide (N2O), methane (CH4), trichlorofluoromethane (CCl3F; commonly referred to as CFC-11) and dichlorodifluoromethane (CCl2F2; commonly referred to as CFC-12), respectively. This data comes from the ESRL/NOAA Global Monitoring Division.
- **Aerosols:** the mean stratospheric aerosol optical depth at 550 nm. This variable is linked to volcanoes, as volcanic eruptions result in new particles being added to the atmosphere, which affect how much of the sun’s energy is reflected back into space. This data is from the Godard Institute for Space Studies at NASA.
- **TSI:** the total solar irradiance (TSI). Due to sunspots and other solar phenomena, the amount of energy that is given off by the sun varies substantially with time. This data is from the SOLARIS-HEPPA project website.
- **MEI:** multivariate El Nino Southern Oscillation index (MEI), a measure of the strength of the El Nino/La Nina-Southern Oscillation (a weather effect in the Pacific Ocean that affects global temperatures). This data comes from the ESRL/NOAA Physical Sciences Division.

Units

- CO2, N2O and CH4 are expressed in **ppmv** (parts per million by volume – i.e., 397 ppmv of CO2 means that CO2 constitutes 397 millionths of the total volume of the atmosphere)
- CFC.11 and CFC.12 are expressed in **ppbv** (parts per billion by volume).
- TSI is expressed in **W/m2** (the rate at which the sun’s energy is deposited per unit area)

First Problem:

We are interested in how changes in these variables affect future temperatures, as well as how well these variables explain temperature changes so far. To do this, first read the dataset climate_change.csv into R.

Data Structure

```
CC<-read.csv("data/climate_change.csv")
```

```
str(CC)
```

```
## 'data.frame': 308 obs. of 11 variables:
## $ Year : int 1983 1983 1983 1983 1983 1983 1983 1983 1984 1984 ...
## $ Month : int 5 6 7 8 9 10 11 12 1 2 ...
## $ MEI : num 2.556 2.167 1.741 1.13 0.428 ...
## $ CO2 : num 346 346 344 342 340 ...
## $ CH4 : num 1639 1634 1633 1631 1648 ...
## $ N2O : num 304 304 304 304 304 ...
## $ CFC.11 : num 191 192 193 194 194 ...
## $ CFC.12 : num 350 352 354 356 357 ...
## $ TSI : num 1366 1366 1366 1366 1366 ...
## $ Aerosols: num 0.0863 0.0794 0.0731 0.0673 0.0619 0.0569 0.0524 0.0486 0.0451 0.0416 ...
## $ Temp : num 0.109 0.118 0.137 0.176 0.149 0.093 0.232 0.078 0.089 0.013 ...
```

Data Summary

```
summary(CC)
```

```
##      Year      Month      MEI      CO2
## Min.   :1983   Min.    : 1.000   Min.   :-1.6350   Min.    :340.2
## 1st Qu.:1989   1st Qu.: 4.000   1st Qu.: -0.3987   1st Qu.:353.0
## Median :1996   Median : 7.000   Median : 0.2375   Median :361.7
## Mean   :1996   Mean   : 6.552   Mean   : 0.2756   Mean   :363.2
## 3rd Qu.:2002   3rd Qu.:10.000   3rd Qu.: 0.8305   3rd Qu.:373.5
## Max.   :2008   Max.   :12.000   Max.    : 3.0010   Max.   :388.5
##      CH4      N2O      CFC.11      CFC.12
## Min.   :1630   Min.    :303.7   Min.    :191.3   Min.    :350.1
## 1st Qu.:1722   1st Qu.:308.1   1st Qu.:246.3   1st Qu.:472.4
## Median :1764   Median :311.5   Median :258.3   Median :528.4
## Mean   :1750   Mean   :312.4   Mean   :252.0   Mean   :497.5
## 3rd Qu.:1787   3rd Qu.:317.0   3rd Qu.:267.0   3rd Qu.:540.5
## Max.   :1814   Max.   :322.2   Max.    :271.5   Max.    :543.8
##      TSI      Aerosols      Temp
## Min.   :1365   Min.    :0.00160   Min.    :-0.2820
## 1st Qu.:1366   1st Qu.:0.00280   1st Qu.: 0.1217
## Median :1366   Median :0.00575   Median : 0.2480
## Mean   :1366   Mean   :0.01666   Mean   : 0.2568
## 3rd Qu.:1366   3rd Qu.:0.01260   3rd Qu.: 0.4073
## Max.   :1367   Max.    :0.14940   Max.    : 0.7390
```

As you can see dataset includes **308** observations and **11** features, namely **Year**, **Month**, **MEI**, **CO2**, **CH4**, **N2O**, **CFC.11**, **CFC.12**, **TSI**, **Aerosols**, **Temp**.

Then, split the data into a training set, consisting of all the observations up to and including 2006, and a testing set consisting of the remaining years (hint: use subset). A training set refers to the data that will be

used to build the model (this is the data we give to the `lm()` function), and a testing set refers to the data we will use to test our predictive ability.

```
trainingset <- subset(CC, Year <= 2006)
testingset <- CC[CC$Year != trainingset$Year,]
```

```
## Warning in CC$Year != trainingset$Year: longer object length is not a
## multiple of shorter object length
```

Next, build a linear regression model to predict the dependent variable Temp, using MEI, CO2, CH4, N2O, CFC.11, CFC.12, TSI, and Aerosols as independent variables (Year and Month should NOT be used in the model). Use the training set to build the model.

```
Model1 <- lm(Temp ~ MEI + CO2 + CH4 + N2O + CFC.11 + CFC.12 + TSI + Aerosols, data = trainingset)
summary(Model1)
```

```
##
## Call:
## lm(formula = Temp ~ MEI + CO2 + CH4 + N2O + CFC.11 + CFC.12 +
##      TSI + Aerosols, data = trainingset)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.25888 -0.05913 -0.00082  0.05649  0.32433
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.246e+02  1.989e+01  -6.265 1.43e-09 ***
## MEI          6.421e-02  6.470e-03   9.923 < 2e-16 ***
## CO2          6.457e-03  2.285e-03   2.826 0.00505 **
## CH4          1.240e-04  5.158e-04   0.240 0.81015
## N2O         -1.653e-02  8.565e-03  -1.930 0.05467 .
## CFC.11       -6.631e-03  1.626e-03  -4.078 5.96e-05 ***
## CFC.12       3.808e-03  1.014e-03   3.757 0.00021 ***
## TSI          9.314e-02  1.475e-02   6.313 1.10e-09 ***
## Aerosols    -1.538e+00  2.133e-01  -7.210 5.41e-12 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.09171 on 275 degrees of freedom
## Multiple R-squared:  0.7509, Adjusted R-squared:  0.7436
## F-statistic: 103.6 on 8 and 275 DF,  p-value: < 2.2e-16
```

Enter the model R2 (the “Multiple R-squared” value): **0.7509**

Problems 1.2:

Which variables are significant in the model?

1. MEI
2. CO2

3. CH4
4. N2O
5. CFC.11
6. CFC.12
7. TSI
8. Aerosols

Answer is: 1, 2,5,6,7,8

Problem 2.1

Current scientific opinion is that nitrous oxide and CFC-11 are greenhouse gases: gases that are able to trap heat from the sun and contribute to the heating of the Earth. However, the regression coefficients of both the N2O and CFC-11 variables are negative, indicating that increasing atmospheric concentrations of either of these two compounds is associated with lower global temperatures.

Which of the following is the simplest correct explanation for this contradiction?

1. Climate scientists are wrong that N2O and CFC-11 are greenhouse gases - this regression analysis constitutes part of a disproof.
2. There is not enough data, so the regression coefficients being estimated are not accurate.
3. All of the gas concentration variables reflect human development - N2O and CFC.11 are correlated with other variables in the data set. All of the gas concentration variables reflect human development - N2O and CFC.11 are correlated with other variables in the data set. - correct

Answer is: 3

Problem 2.2

Compute the correlations between all the variables in the training set. Which of the following independent variables is N2O highly correlated with (absolute correlation greater than 0.7)? Select all that apply.

```
cor(trainingset)
```

##	Year	Month	MEI	CO2	CH4
## Year	1.00000000	-0.0279419602	-0.0369876842	0.98274939	0.91565945
## Month	-0.02794196	1.0000000000	0.0008846905	-0.10673246	0.01856866
## MEI	-0.03698768	0.0008846905	1.0000000000	-0.04114717	-0.03341930
## CO2	0.98274939	-0.1067324607	-0.0411471651	1.00000000	0.87727963
## CH4	0.91565945	0.0185686624	-0.0334193014	0.87727963	1.00000000
## N2O	0.99384523	0.0136315303	-0.0508197755	0.97671982	0.89983864
## CFC.11	0.56910643	-0.0131112236	0.0690004387	0.51405975	0.77990402
## CFC.12	0.89701166	0.0006751102	0.0082855443	0.85268963	0.96361625
## TSI	0.17030201	-0.0346061935	-0.1544919227	0.17742893	0.24552844
## Aerosols	-0.34524670	0.0148895406	0.3402377871	-0.35615480	-0.26780919
## Temp	0.78679714	-0.0998567411	0.1724707512	0.78852921	0.70325502
##	N2O	CFC.11	CFC.12	TSI	Aerosols
## Year	0.99384523	0.56910643	0.8970116635	0.17030201	-0.34524670
## Month	0.01363153	-0.01311122	0.0006751102	-0.03460619	0.01488954
## MEI	-0.05081978	0.06900044	0.0082855443	-0.15449192	0.34023779

```
## CO2      0.97671982  0.51405975  0.8526896272  0.17742893 -0.35615480
## CH4      0.89983864  0.77990402  0.9636162478  0.24552844 -0.26780919
## N2O      1.00000000  0.52247732  0.8679307757  0.19975668 -0.33705457
## CFC.11   0.52247732  1.00000000  0.8689851828  0.27204596 -0.04392120
## CFC.12   0.86793078  0.86898518  1.0000000000  0.25530281 -0.22513124
## TSI      0.19975668  0.27204596  0.2553028138  1.00000000  0.05211651
## Aerosols -0.33705457 -0.04392120 -0.2251312440  0.05211651  1.00000000
## Temp     0.77863893  0.40771029  0.6875575483  0.24338269 -0.38491375
##          Temp
## Year      0.78679714
## Month     -0.09985674
## MEI       0.17247075
## CO2       0.78852921
## CH4       0.70325502
## N2O       0.77863893
## CFC.11    0.40771029
## CFC.12    0.68755755
## TSI       0.24338269
## Aerosols  -0.38491375
## Temp      1.00000000
```

Answer is: CO2, CH4, CFC.12

Which of the following independent variables is CFC.11 highly correlated with? Select all that apply.

Answer is CH4,CFC.12

Problem 3 - Simplifying the Model

(2 points possible) Given that the correlations are so high, let us focus on the N2O variable and build a model with only MEI, TSI, Aerosols and N2O as independent variables. Remember to use the training set to build the model.

```
Model2 <- lm(Temp ~ MEI + TSI + Aerosols + N2O, data = trainingset)
summary(Model2)
```

```
##
## Call:
## lm(formula = Temp ~ MEI + TSI + Aerosols + N2O, data = trainingset)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.27916 -0.05975 -0.00595  0.05672  0.34195
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.162e+02  2.022e+01  -5.747 2.37e-08 ***
## MEI          6.419e-02  6.652e-03   9.649 < 2e-16 ***
## TSI          7.949e-02  1.487e-02   5.344 1.89e-07 ***
## Aerosols    -1.702e+00  2.180e-01  -7.806 1.19e-13 ***
## N2O          2.532e-02  1.311e-03  19.307 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Residual standard error: 0.09547 on 279 degrees of freedom
## Multiple R-squared:  0.7261, Adjusted R-squared:  0.7222
## F-statistic: 184.9 on 4 and 279 DF,  p-value: < 2.2e-16
```

Enter the coefficient of N2O in this reduced model: **2.532e-02**

How does this compare to the coefficient in the previous model with all of the variables? ** The R-squared is decreased**

Enter the model R2: **0.7261**

Problem 4 - Automatically Building the Model

We have many variables in this problem, and as we have seen above, dropping some from the model does not decrease model quality. R provides a function, `step`, that will automate the procedure of trying different combinations of variables to find a good compromise of model simplicity and R2. This trade-off is formalized by the Akaike information criterion (AIC) - it can be informally thought of as the quality of the model with a penalty for the number of variables in the model.

The `step` function has one argument - the name of the initial model. It returns a simplified model. Use the `step` function in R to derive a new model, with the full model as the initial model (HINT: If your initial full model was called “climateLM”, you could create a new model with the `step` function by typing `step(climateLM)`. Be sure to save your new model to a variable name so that you can look at the summary. For more information about the `step` function, type `?step` in your R console.)

```
Model_UsigStep <- step(Model1)
```

```
## Start:  AIC=-1348.16
## Temp ~ MEI + CO2 + CH4 + N2O + CFC.11 + CFC.12 + TSI + Aerosols
##
##           Df Sum of Sq    RSS    AIC
## - CH4      1  0.00049 2.3135 -1350.1
## <none>      0          2.3130 -1348.2
## - N2O      1  0.03132 2.3443 -1346.3
## - CO2      1  0.06719 2.3802 -1342.0
## - CFC.12   1  0.11874 2.4318 -1335.9
## - CFC.11   1  0.13986 2.4529 -1333.5
## - TSI      1  0.33516 2.6482 -1311.7
## - Aerosols 1  0.43727 2.7503 -1301.0
## - MEI      1  0.82823 3.1412 -1263.2
##
## Step:  AIC=-1350.1
## Temp ~ MEI + CO2 + N2O + CFC.11 + CFC.12 + TSI + Aerosols
##
##           Df Sum of Sq    RSS    AIC
## <none>      0          2.3135 -1350.1
## - N2O      1  0.03133 2.3448 -1348.3
## - CO2      1  0.06672 2.3802 -1344.0
## - CFC.12   1  0.13023 2.4437 -1336.5
## - CFC.11   1  0.13938 2.4529 -1335.5
## - TSI      1  0.33500 2.6485 -1313.7
## - Aerosols 1  0.43987 2.7534 -1302.7
## - MEI      1  0.83118 3.1447 -1264.9
```

```
summary(Model_UsigStep)
```

```
##
## Call:
## lm(formula = Temp ~ MEI + CO2 + N2O + CFC.11 + CFC.12 + TSI +
##     Aerosols, data = trainingset)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.25770 -0.05994 -0.00104  0.05588  0.32203
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.245e+02  1.985e+01  -6.273 1.37e-09 ***
## MEI          6.407e-02  6.434e-03   9.958 < 2e-16 ***
## CO2          6.402e-03  2.269e-03   2.821 0.005129 **
## N2O         -1.602e-02  8.287e-03  -1.933 0.054234 .
## CFC.11       -6.609e-03  1.621e-03  -4.078 5.95e-05 ***
## CFC.12       3.868e-03  9.812e-04   3.942 0.000103 ***
## TSI          9.312e-02  1.473e-02   6.322 1.04e-09 ***
## Aerosols    -1.540e+00  2.126e-01  -7.244 4.36e-12 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.09155 on 276 degrees of freedom
## Multiple R-squared:  0.7508, Adjusted R-squared:  0.7445
## F-statistic: 118.8 on 7 and 276 DF,  p-value: < 2.2e-16
```

Enter the R2 value of the model produced by the step function:

NOTE:It is interesting to note that the step function does not address the collinearity of the variables, except that adding highly correlated variables will not improve the R2 significantly. The consequence of this is that the step function will not necessarily produce a very interpretable model - just a model that has balanced quality and simplicity for a particular weighting of quality and simplicity (AIC).